

Mapeo Objeto Relacional

Merging Relational and Object Models

- ▶ Object-oriented models support interesting data types --- not just flat files.
 - ▶ Maps, multimedia, etc.
- ▶ The relational model supports very-high-level queries.
- ▶ Object-relational databases are an attempt to get the best of both.

Complex Data Types

- ▶ Motivation:
 - ▶ Permit non-atomic domains (atomic \equiv indivisible)
 - ▶ Example of non-atomic domain: set of integers, or set of tuples
 - ▶ Allows more intuitive modeling for applications with complex data
- ▶ Intuitive definition:
 - ▶ Retains mathematical foundation of relational model
 - ▶ Violates first normal form.

Example of a Nested Relation

- ▶ Example: library information system
- ▶ Each book has
 - ▶ title,
 - ▶ a list (array) of authors,
 - ▶ Publisher, with subfields *name* and *branch*, and
 - ▶ a set of keywords
- ▶ Non-1NF relation *books*

<i>title</i>	<i>author_array</i>	<i>publisher</i>	<i>keyword_set</i>
		(<i>name</i> , <i>branch</i>)	
Compilers	[Smith, Jones]	(McGraw-Hill, NewYork)	{parsing, analysis}
Networks	[Jones, Frick]	(Oxford, London)	{Internet, Web}

Complex Types and SQL

- ▶ Extensions introduced in SQL:1999 to support complex types:
 - ▶ Collection and large object types
 - ▶ Nested relations are an example of collection types
 - ▶ Structured types
 - ▶ Nested record structures like composite attributes
 - ▶ Inheritance
 - ▶ Object orientation
 - ▶ Including object identifiers and references
- ▶ Not fully implemented in any database system currently
 - ▶ But some features are present in each of the major commercial database systems
 - ▶ Read the manual of your database system to see what it supports

User Defined Types

- ▶ A *user-defined type*, or UDT, is essentially a class definition, with a structure and methods.
- ▶ Two uses:
 1. As a *rowtype*, that is, the type of a relation.
 2. As the type of an attribute of a relation.

Structured Types and Inheritance in SQL

- ▶ **Structured types** (a.k.a. **user-defined types**) can be declared and used in SQL

```
create type Name as
  (firstname      varchar(20),
   lastname      varchar(20))
final
```

```
create type Address as
  (street        varchar(20),
   city          varchar(20),
   zipcode       varchar(20))
```

not final

- ▶ Note: **final** and **not final** indicate whether subtypes can be created
- ▶ Structured types can be used to create tables with composite attributes

```
create table person (
  name      Name,
  address   Address,
  dateOfBirth date)
```

- ▶ Dot notation used to reference components: *name.firstname*

Structured Types (cont.)

- ▶ **User-defined row types**

```
create type PersonType as (  
    name Name,  
    address Address,  
    dateOfBirth date)  
not final
```

- ▶ Can then create a table whose rows are a user-defined type
create table *customer* of *CustomerType*

- ▶ Alternative using **unnamed row types**.

```
create table person_r(  
    name    row(firstname varchar(20),  
                lastname  varchar(20)),  
    address row(street    varchar(20),  
                city      varchar(20),  
                zipcode   varchar(20)),  
    dateOfBirth date)
```


Constructor Functions

- ▶ **Constructor functions** are used to create values of structured types
- ▶ E.g.
create function *Name*(*firstname* varchar(20), *lastname* varchar(20))
returns *Name*
begin
 set self.*firstname* = *firstname*;
 set self.*lastname* = *lastname*;
end
- ▶ To create a value of type *Name*, we use
 new *Name*('John', 'Smith')
- ▶ Normally used in insert statements
insert into *Person* values
 (new *Name*('John', 'Smith'),
 new *Address*('20 Main St', 'New York', '11001'),
 date '1960-8-22');

Type Inheritance

- ▶ Suppose that we have the following type definition for people:

```
create type Person  
  (name varchar(20),  
   address varchar(20))
```
- ▶ Using inheritance to define the student and teacher types

```
create type Student  
  under Person  
  (degree      varchar(20),  
   department varchar(20))  
create type Teacher  
  under Person  
  (salary      integer,  
   department varchar(20))
```
- ▶ Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration

Multiple Type Inheritance

- ▶ SQL:1999 and SQL:2003 do not support multiple inheritance
- ▶ If our type system supports multiple inheritance, we can define a type for teaching assistant as follows:
create type *Teaching Assistant*
under *Student, Teacher*
- ▶ To avoid a conflict between the two occurrences of *department* we can rename them

create type *Teaching Assistant*
under
Student with (department as student_dept),
Teacher with (department as teacher_dept)
- ▶ Each value must have a **most-specific type**

Array and Multiset Types in SQL

- Example of array and multiset declaration:

```
create type Publisher as
  (name          varchar(20),
   branch       varchar(20));
create type Book as
  (title        varchar(20),
   author_array varchar(20) array [10],
   pub_date     date,
   publisher    Publisher,
   keyword-set  varchar(20) multiset);
create table books of Book;
```

Creation of Collection Values

- ▶ Array construction

```
array ['Silberschatz', `Korth`, `Sudarshan']
```

- ▶ Multisets

```
multiset ['computer', 'database', 'SQL']
```

- ▶ To create a tuple of the type defined by the books relation:

```
('Compilers', array[ `Smith`, `Jones`,  
                    new Publisher ( `McGraw-Hill`, `New York` ),  
multiset [ `parsing`, `analysis` ] )
```

- ▶ To insert the preceding tuple into the relation books

```
insert into books  
values
```

```
( 'Compilers', array[ `Smith`, `Jones`,  
                    new Publisher ( `McGraw-Hill`, `New York` ),  
multiset [ `parsing`, `analysis` ] );
```

Unnesting

- ▶ The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**.

- ▶ E.g.

```
select title, A as author, publisher.name as pub_name,  
       publisher.branch as pub_branch, K.keyword  
from books as B, unnest(B.author_array ) as A (author ),  
       unnest (B.keyword_set ) as K (keyword )
```

- ▶ Result relation *flat_books*

<i>title</i>	<i>author</i>	<i>pub_name</i>	<i>pub_branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web

Querying Collection-Valued Attributes

- ▶ To find all books that have the word “database” as a keyword,

```
select title  
from books  
where 'database' in (unnest(keyword-set ))
```
- ▶ We can access individual elements of an array by using indices
 - ▶ E.g.: If we know that a particular book has three authors, we could write:

```
select author_array[1], author_array[2], author_array[3]  
from books  
where title = `Database System Concepts`
```
- ▶ To get a relation containing pairs of the form “title, author_name” for each book and each author of the book

```
select B.title, A.author  
from books as B, unnest (B.author_array) as A (author )
```
- ▶ To retain ordering information we add a **with ordinality** clause

```
select B.title, A.author, A.position  
from books as B, unnest (B.author_array) with ordinality as  
A (author, position )
```

Nesting

- ▶ **Nesting** is the opposite of unnesting, creating a collection-valued attribute
- ▶ Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset

- ▶ To nest the *flat_books* relation on the attribute *keyword*:

```
select title, author, Publisher (pub_name, pub_branch ) as publisher,  
       collect (keyword) as keyword_set  
from flat_books  
groupby title, author, publisher
```

- ▶ To nest on both authors and keywords:

```
select title, collect (author ) as author_set,  
       Publisher (pub_name, pub_branch) as publisher,  
       collect (keyword ) as keyword_set  
from flat_books  
group by title, publisher
```


Nesting (Cont.)

- ▶ Another approach to creating nested relations is to use subqueries in the **select** clause,
- ▶ **select** *title*,
 array (**select** *author*
 from *authors* **as** *A*
 where *A.title* = *B.title*
 order by *A.position*) **as** *author_array*,
 Publisher (*pub-name*, *pub-branch*) **as** *publisher*,
 multiset (**select** *keyword*
 from *keywords* **as** *K*
 where *K.title* = *B.title*) **as** *keyword_set*
from *books4* **as** *B*

Storing Nested Relations

- ▶ Oracle doesn't really store each nested table as a separate relation --- it just makes it look that way.
- ▶ Rather, there is one relation R in which all the tuples of all the nested tables for one attribute A are stored.
- ▶ Declare in CREATE TABLE by:
 NESTED TABLE A STORE AS R

Example: Storing Nested Tables

```
CREATE TABLE Manfs (  
    name  CHAR(30),  
    addr   CHAR(50),  
    beers  beerTableType  
)  
NESTED TABLE beers STORE AS BeerTable;
```

References

- ▶ If T is a type, then $\text{REF } T$ is the type of a reference to T , that is, a pointer to an object of type T .
- ▶ Often called an “object ID” in OO systems.
- ▶ Unlike object ID’s, a REF is visible, although it is gibberish.

Object-Identity and Reference Types

- ▶ Define a type *Department* with a field *name* and a field *head* which is a reference to the type *Person*, with table *people* as scope:

```
create type Department (  
    name varchar (20),  
    head ref (Person) scope people)
```

- ▶ We can then create a table *departments* as follows

```
create table departments of Department
```

- ▶ We can omit the declaration **scope people** from the type declaration and instead make an addition to the **create table** statement:

```
create table departments of Department  
    (head with options scope people)
```

- ▶ Referenced table must have an attribute that stores the identifier, called the **self-referential attribute**

```
create table people of Person  
ref is person_id system generated;
```

Initializing Reference-Typed Values

- To create a tuple with a reference value, we can first create the tuple with a null reference and then set the reference separately:

```
insert into departments
  values (`CS`, null)
update departments
  set head = (select p.person_id
              from people as p
              where name = `John`)
  where name = `CS`
```

Object Identifiers Using Reference Types

► Reference type

- Create unique system-generated object identifiers

► Examples:

- `REF IS SYSTEM GENERATED`
- `REF IS <OID_ATTRIBUTE>`
`<VALUE_GENERATION_METHOD> ;`

User Generated Identifiers

- ▶ The type of the object-identifier must be specified as part of the type definition of the referenced table, and
- ▶ The table definition must specify that the reference is user generated

```
create type Person
  (name varchar(20)
   address varchar(20))
  ref using varchar(20)
create table people of Person
  ref is person_id user generated
```

- ▶ When creating a tuple, we must provide a unique value for the identifier:

```
insert into people (person_id, name, address ) values
  ('01284567', 'John', '23 Coyote Run')
```
- ▶ We can then use the identifier value when inserting a tuple into *departments*
 - ▶ Avoids need for a separate query to retrieve the identifier:

```
insert into departments
values( 'CS', '02184567')
```


User Generated Identifiers (Cont.)

- ▶ Can use an existing primary key value as the identifier:

```
create type Person
  (name varchar (20) primary key,
   address varchar(20))
ref from (name)
create table people of Person
  ref is person_id derived
```

- ▶ When inserting a tuple for *departments*, we can then use

```
insert into departments
  values(`CS`, `John`)
```

Path Expressions

- ▶ Find the names and addresses of the heads of all departments:

```
select head ->name, head ->address  
from departments
```

- ▶ An expression such as “head->name” is called a **path expression**
- ▶ Path expressions help avoid explicit joins
 - ▶ If department head were not a reference, a join of *departments* with *people* would be required to get at the address
 - ▶ Makes expressing the query much easier for the user

Implementing O-R Features

- ▶ Similar to how E-R features are mapped onto relation schemas
- ▶ Subtable implementation
 - ▶ Each table stores primary key and those attributes defined in that tableor,
 - ▶ Each table stores both locally defined and inherited attributes

Presentación

- ▶ Esta presentación fue armada utilizando, además de material propio, material provisto por los siguientes autores:
- ▶ Siblberschat, Korth, Sudarshan - Database Systems Concepts, 6th Ed., Mc Graw Hill, 2010
- ▶ García Molina/Ullman/Widom - Database Systems: The Complete Book, 2nd Ed., Prentice Hall, 2009
- ▶ Elmasri/Navathe - Fundamentals of Database Systems, 6th Ed., Addison Wesley, 2011